

Stop and Start feature

- This feature's purpose is to help us save costs
- A feature designed for AKS clusters that don't sustain workloads that need to run non-stop
- Stopping a cluster results in:
 - The control plane getting terminated -> API server will not be responsive. After starting the cluster, the IP address of the API server may change
 - Nodes will be stopped -> full compute savings, because the VM instances are deprovisioned. If we had stopped the VMSS/VMs, that would've meant they were deallocated, so we would've still been charged for the OS
- A cluster should be stopped for up to 1 year (12 months). After this, its state is unrecoverable
- It is recommended to wait about 15 minutes before starting the cluster after you stopped it
- You can separately stop node pools, but only User node pools, not System node pools
- Stopping process drains all nodes, so the standalone pods will be deleted and not recreated after the cluster is started. A stand-alone pod is one that is not managed by a higher-level Kubernetes object such as Deployment, DaemonSet, StatefulSet, and so on